



SIMATS
ENGINEERING



SIMATS
Saveetha Institute of Medical And Technical Sciences
(Declared as Deemed to be University under Section 3 of UGC Act 1956)

**A MULTI-AGENT INTELLIGENT FIREWALL FOR PROMPT
INJECTION, HALLUCINATION DETECTION, AND SECURE
ENTERPRISE LLM APPLICATIONS**

A CAPSTONE PROJECT REPORT

*A Capstone Project Report submitted in partial fulfilment of the requirements for
the Course of*

CSA6502 – GENERATIVE AI AND LARGE SCALE MODELS

to the award of the degree of

BACHELOR OF ENGINEERING / BACHELOR OF TECHNOLOGY

IN

COMPUTER SCIENCE ENGINEERING - AI

Submitted by

S. Muskan(192472182)

Under the Supervision of

Dr. Nithisha J

SIMATS ENGINEERING

Saveetha Institute of Medical and Technical Sciences

Chennai-602105

SEPTEMBER 2026

SIMATS ENGINEERING

Saveetha Institute of Medical and Technical Sciences

Chennai-602105

BONAFIDE CERTIFICATE

This is to certify that the Capstone Project entitled “**A Multi-Agent Intelligent Firewall for Prompt Injection, Hallucination Detection, and Secure Enterprise LLM Applications**” has been carried out by **S. Muskan(192472182)** under the supervision of **Dr. Nithisha J** and is submitted in partial fulfilment of the requirements for the current semester of the **B.E.CSE(AI)** program at Saveetha Institute of Medical and Technical Sciences, Chennai

COURSE COORDINATOR	COURSE FACULTY
Dr. Chandravathi C	Dr. Nithisha J
Program Director	Professor
CSE	CSE
SIMATS Engineering,	SIMATS Engineering,
SIMATS, Chennai – 602105.	SIMATS, Chennai – 602105.

SIMATS ENGINEERING

Saveetha Institute of Medical and Technical Sciences

Chennai-602105

DECLARATION

I, **S. Muskan(192472182)** of the **B.E.CSE(AI), SIMATS Engineering**, Saveetha Institute of Medical and Technical Sciences, Chennai, hereby declare that the Capstone Project Work entitled **“A Multi-Agent Intelligent Firewall for Prompt Injection, Hallucination Detection, and Secure Enterprise LLM Applications”** is the result of our my own bonafide efforts. To the best of my knowledge, the work presented herein is original, accurate, and has been carried out in accordance with the principles of engineering ethics and academic integrity. All sources, references, data, and other materials used in the project have been appropriately acknowledged. I further declare that the data, results, analysis, and findings presented in this report have not been fabricated, falsified, or manipulated. Any external tools, software, datasets, computational resources, or AI-assisted tools used during the project have been appropriately disclosed. We confirm that all applicable ethical, academic, institutional, and professional requirements have been duly followed throughout the planning, execution, analysis, and documentation of the project.

Place:

Date:

Signature of the Students with Names

S.Muskan

INDIVIDUAL CONTRIBUTION STATEMENT

Student	Specific Responsibility	Contribution Area	Approx. Contribution
B. Vani / 192472348	Enterprise AI security requirements and coordination	Problem definition, architecture and documentation	25%
B. Paneesha / 192472206	Prompt security and firewall workflow	Prompt injection/jailbreak protection and analysis	25%
R. Mouni / 192472283	Hallucination and sensitive-data protection	Response verification, RAG and data protection	25%
S. Muskan / 192472182	Evaluation, logging and reporting	Risk scoring, audit logging, testing and presentation	25%
TOTAL			100%

ABSTRACT

Large Language Models are transforming enterprise applications through intelligent automation, but their adoption introduces new security and reliability challenges. Enterprise AI systems can be exposed to prompt injection and jailbreak attacks, hallucinated responses can reduce trust in generated information, and sensitive organizational data may be disclosed through insecure AI interactions. Traditional network-oriented firewalls and manual monitoring approaches do not adequately address these language-model-specific risks.

This project proposes a Multi-Agent Intelligent Firewall for Prompt Injection, Hallucination Detection, and Secure Enterprise LLM Applications. The framework authenticates users, applies role-based access control, routes requests through an LLM gateway, analyzes prompts for malicious instructions and jailbreak attempts, verifies generated responses, detects sensitive information, assigns risk scores, and maintains audit logs.

The framework is organized into four major modules: User Interface, Authentication & Access Control and LLM Gateway; Multi-Agent Firewall & Prompt Injection Detection Agent; Hallucination Detection Agent & Sensitive Data Protection Agent; and Risk Scoring Engine, Audit Logging, Database, Admin Dashboard & Model Evaluation. RAG and semantic similarity are identified for response verification.

The expected outcome is a secure, reliable and scalable framework for enterprise Generative AI applications. The presentation identifies future extensions including multimodal AI, adaptive learning, cloud deployment, SIEM integration, autonomous security decisions and zero-trust AI architecture.

Keywords: Large Language Models (LLMs), Multi-Agent Systems, Prompt Injection Detection, Jailbreak Detection, Hallucination Detection, Sensitive Data Protection, Risk Scoring, Retrieval-Augmented Generation (RAG), Enterprise AI Security, Audit Logging.

TABLE OF CONTENTS

Sl. No	Title	Page No
i	CERTIFICATE FROM THE SUPERVISOR	ii
ii	DECLARATION BY THE CANDIDATE	iii
Iii	INDIVIDUAL CONTRIBUTION STATEMENT	iv
iv	ABSTRACT	v
v	LIST OF FIGURES	vi
vi	LIST OF TABLES	vii
vii	LIST OF ABBREVIATIONS	viii
1	CHAPTER 1: Introduction	10
2	CHAPTER 2: Literature Review	13
3	CHAPTER 3: Engineering Design & Trade-offs, Sustainability, Ethics, Safety, Risk & Security	15
4	CHAPTER 4: Implementation, Testing & Results, Project Management	16
5	CHAPTER 5: Conclusion & Reflection	18
	References	19
	Appendices	20-251

LIST OF TABLES

Table No.	Table Name	Page No.
Table 2.2	Comparison of Existing Healthcare Management Approaches	6
Table 3.1	Stakeholder Requirements	9
Table 3.2	Functional Requirements	10
Table 3.3	Non-Functional Requirements	10
Table 3.4	Design Specifications	11
Table 3.5	Alternative Solution Evaluation	12
Table 3.6	Trade-off Analysis	13
Table 3.7	Sustainability, Ethics, Safety and Security Analysis	14
Table 3.8	Risk Register	14
Table 4.1	Software Tools and Resources Used	16
Table 4.2	Test Plan and Verification	17
Table 4.3	Test Results	18
Table 4.4	Objective Achievement	19
Table 4.5	Strengths and Limitations	19
Table 4.6	Project Timeline	20
Table 4.7	Project Roles and Contributions	20
Table 4.8	Project Budget	20

List of Figures

Fig No.	Table Name	Page No.
Fig. 3.1	Multi - Agent Intelligent Firewall Architecture	15
Fig. 4.1	Secure Chat Interface Showing Prompt Security Detection	24
Fig. 4.2	Admin Dashboard Showing Model Evaluation and Audit Logs	24
Fig. 4.3	Request Logs Showing Security Actions and Risk Scores	25

LIST OF ABBREVIATIONS

Abbreviation	Description
AI	Artificial Intelligence
LLM	Large Language Model
RAG	Retrieval-Augmented Generation
UI	User Interface
RBAC	Role-Based Access Control
SIEM	Security Information and Event Management

CHAPTER 1

INTRODUCTION AND ENGINEERING PROBLEM

1.1 Background

Large Language Models are transforming enterprise applications through intelligent automation. However, enterprise AI systems are increasingly vulnerable to prompt injection and jailbreak attacks. Hallucinated responses can reduce reliability, while insecure AI interactions can expose sensitive enterprise data. The proposed firewall introduces an intelligent security layer around enterprise LLM applications.

1.2 Need for the Project

- Enterprise LLM applications are vulnerable to prompt injection attacks.
- AI models may generate hallucinated or misleading responses.
- Sensitive organizational information can be exposed during AI interactions.
- Existing security mechanisms provide limited AI-specific protection.
- Manual monitoring is inefficient and time-consuming.
- Organizations require an intelligent framework to secure enterprise LLMs.

1.3 Problem Statement

The engineering problem is to design an intelligent multi-agent security framework that inspects enterprise AI prompts before model execution, detects prompt injection and jailbreak attempts, verifies generated responses for hallucination, protects sensitive information, assigns risk scores and maintains auditable security records.

1.4 Aim

To develop a secure multi-agent intelligent firewall that improves the safety, reliability, trustworthiness and monitoring of enterprise Large Language Model applications.

The engineering problem is to design an intelligent multi-agent security framework that inspects enterprise AI prompts before model execution, detects prompt injection and jailbreak attempts, verifies generated responses for hallucination, protects sensitive information, assigns risk scores and maintains auditable security records.

1.5 Objectives

- Develop a secure multi-agent firewall for enterprise LLM applications.
- Verify AI-generated responses to reduce hallucinated content.
- Protect sensitive enterprise information during AI interactions.
- Generate audit logs for security monitoring and compliance.
- Improve the reliability and trustworthiness of enterprise AI systems.
- Detect prompt injection and jailbreak attacks before model execution.
- Verify AI-generated responses to reduce hallucinated content.
- Protect sensitive enterprise information during AI interactions.

- Generate audit logs for security monitoring and compliance.
- Improve the reliability and trustworthiness of enterprise AI systems.

1.6 Scope

The project covers secure user access, prompt inspection, threat detection, response verification, sensitive-data protection, risk scoring, audit logging, database support, administrative monitoring and model evaluation. The engineering problem is to design an intelligent multi-agent security framework that inspects enterprise AI prompts before model execution, detects prompt injection and jailbreak attempts, verifies generated responses for hallucination, protects sensitive information, assigns risk scores and maintains auditable security records. The source presentation does not provide completed numerical experimental results, so quantitative claims are intentionally not added.

1.7 Expected Engineering Outcome

A unified security framework capable of detecting malicious prompts, reducing hallucinated responses, protecting sensitive enterprise information and providing continuous monitoring.

The engineering problem is to design an intelligent multi-agent security framework that inspects enterprise AI prompts before model execution, detects prompt injection and jailbreak attempts, verifies generated responses for hallucination, protects sensitive information, assigns risk scores and maintains auditable security records.

CHAPTER 2

LITERATURE REVIEW AND EXISTING SOLUTIONS

2.1 Review Method

The literature survey focuses on prompt injection, jailbreak attacks, hallucination detection and autonomous red teaming of Large Language Models.

2.2 Literature Survey

S.No.	Paper / Work	Authors	Methodology	Relevance
1	Goal-Guided Generative Prompt Injection Attack on Large Language Models (2025)	Chong Zhang et al.	Goal-oriented adversarial injection generation using an auxiliary model.	Shows targeted prompt-injection vulnerabilities.
2	Jailbreak Attacks on Large Language Models and Possible Defenses (2024)	Sadaf Surur Ahmed; J. Angel Arul Jothi	Review of jailbreak attacks and defenses.	Provides threat/defense context.
3	MixHD: Hallucination Detection (2025)	Chuang Li et al.	Combines internal states and output probabilities.	Motivates multi-signal verification.
4	RedAgent: Context-Aware Red Teaming (2026)	Huiyu Xu et al.	Autonomous multi-agent jailbreak red teaming.	Supports multi-agent security evaluation.

2.3 Existing Solution

- Traditional firewalls inspect network traffic rather than natural-language prompts.
- Sensitive enterprise information remains vulnerable.
- Current frameworks provide limited AI-specific protection.
- Manual monitoring increases operational complexity.
- Existing systems may not effectively detect prompt injection.
- AI responses are rarely verified for factual correctness.
- Sensitive enterprise information remains vulnerable.
- Current frameworks provide limited AI-specific protection.
- Manual monitoring increases operational complexity.

2.4 Engineering Gap

The identified gap is limited integration of prompt-level threat detection, jailbreak detection, hallucination verification, sensitive-data protection, risk assessment and audit monitoring in one enterprise AI workflow.

2.5 Proposed Contribution

The proposed contribution is a multi-agent intelligent firewall in which specialized agents collaborate around the enterprise LLM. The identified gap is limited integration of prompt-level threat detection, jailbreak detection, hallucination verification, sensitive-data protection, risk assessment and audit monitoring in one enterprise AI workflow. Prompts are analyzed before model execution, responses are verified before delivery, sensitive information is protected and security events are logged and scored.

CHAPTER 3

ENGINEERING DESIGN & TRADE-OFFS, SUSTAINABILITY, ETHICS, SAFETY, RISK & SECURITY

3.1 Stakeholder Requirements

Stakeholder	Requirement
Enterprise AI User	Secure and understandable LLM access.
Security Administrator	Real-time monitoring, risk scores and audit logs.
Enterprise Organization	Protection of confidential information and reliable AI outputs.
Developer / Evaluator	Clear modules and measurable evaluation.

3.2 Functional and Non-Functional Requirements

Type	Requirement	Target
Functional	Authentication and access control	Verify identity and enforce permissions.
Functional	Prompt injection detection	Inspect every prompt before execution.
Functional	Jailbreak detection	Identify attempts to bypass safety policies.
Functional	Hallucination verification	Check generated responses.
Functional	Sensitive-data protection	Mask or block confidential information.
Functional	Risk scoring and audit logging	Record decisions and security scores.
Non-Functional	Reliability	Consistent decisions for equivalent inputs.
Non-Functional	Usability	Clear monitoring views.
Non-Functional	Security	Protected access and logs.

3.3 System Architecture




Fig 3.1 - Multi - Agent Intelligent Firewall Architecture

3.4 Alternative Solutions & Trade-offs

Approach	Benefit	Limitation	Decision
Network firewall only	Established network security	Does not understand prompt-level attacks	Insufficient alone
Manual monitoring	Human judgment	Slow and difficult at scale	Not primary approach
Single security model	Simpler	Less specialization	Not selected
Multi-agent firewall	Specialized collaborative analysis	Higher integration complexity	Selected

3.5 Sustainability, Ethics, Safety & Security

Domain	Consideration	Impact
Sustainability	Software-based security architecture	Avoids additional physical security hardware.
Ethics	Transparent reporting and responsible AI use	Improves accountability.
Safety	Prompt and response screening	Reduces unsafe interactions.

Security	Authentication, access control, data protection and audit logs	Reduces unauthorized access and disclosure.
----------	--	---

3.6 Risk Register

Risk	Likelihood	Severity	Mitigation
Prompt detection error	Medium	High	Multi-agent evaluation and continuous testing.
Hallucination verification failure	Medium	High	Trusted knowledge and response verification.
Sensitive data leakage	Medium	High	Detection, masking/blocking and access control.
Audit log compromise	Low/Medium	High	Secure storage and restricted access.
Agent integration error	Medium	Medium	Module-level testing and controlled integration.

CHAPTER 4

IMPLEMENTATION, TESTING & RESULTS, PROJECT MANAGEMENT

4.1 Development Approach

The system is designed as a modular workflow: authentication and access control, LLM Gateway, multi-agent prompt analysis, enterprise LLM execution, response verification, sensitive-data protection, risk scoring and audit logging.

4.2 Implementation Modules

- Secure user interface and authentication.
- Role-based access control.
- LLM Gateway.
- Prompt injection detection agent.
- Jailbreak detection agent.
- Hallucination detection using RAG and semantic similarity concepts.
- Sensitive data protection agent.
- Risk scoring engine.
- Audit logging and database.
- Admin dashboard and model evaluation.

4.3 Test Plan & Verification

Test Area	Test Method	Acceptance Criterion
Authentication	Valid/invalid access attempts	Only authorized users gain access.
Access Control	Test different roles	Permissions match roles.
Prompt Injection	Benign and malicious-style prompts	Unsafe prompts detected/blocked/sanitized.
Jailbreak	Bypass-style prompts	Attempts identified.
Hallucination	Compare response with trusted knowledge	Unsupported content flagged/reduced.
Sensitive Data	Responses containing confidential patterns	Data masked or blocked.
Risk Scoring	Interactions with varied threat levels	Scores reflect security decisions.
Audit Logging	Representative interactions	Prompts, responses and decisions recorded.
Dashboard	Review monitoring view	Events are traceable and understandable.

4.4 Expected Results

The source presentation describes expected outputs rather than measured experimental results. Expected behavior includes secure authentication, detection/blocking of prompt injection and jailbreak attempts, reduction of hallucinated responses through verification, protection of sensitive enterprise data, and continuous monitoring through risk scores and audit logs.

Expected Output	Expected Observation
Secure authentication	Unauthorized users are prevented from protected functions.
Prompt protection	Malicious prompt patterns are detected before execution.
Jailbreak protection	Attempts to bypass safety policies are flagged or blocked.
Response verification	Potential hallucinations are identified.
Sensitive data protection	Confidential information is masked or blocked.
Risk monitoring	Interactions receive security-oriented risk scores.
Auditability	Security events and decisions are logged.

4.5 Data Analysis & Engineering Judgment

Evaluation should consider detection effectiveness, false positives and negatives, response verification quality, sensitive-data protection, risk-score consistency, logging completeness and system performance. Numerical accuracy, latency or detection-rate claims should be inserted only after actual implementation and testing.

4.6 Strengths & Limitations

- Integrated protection across prompts and generated responses.
- Specialized agents address different enterprise AI security risks.
- Risk scoring and audit logging improve monitoring and accountability.
- Architecture can be extended to enterprise-scale environments.
- Multi-agent coordination increases complexity.
- Detection quality depends on model/agent capability and evaluation data.

4.7 Project Planning & Roles

Milestone	Activity
1	Problem identification and requirements
2	Literature survey and threat analysis
3	Architecture and module design
4	Firewall and detection implementation
5	Response verification and data protection
6	Risk scoring, logging and dashboard
7	Testing and evaluation
Student	Primary Responsibility
S.Muskan	Prompt security and injection/jailbreak protection

CHAPTER 5

CONCLUSION, FUTURE WORK AND REFLECTION

5.1 Conclusion

The proposed multi-agent intelligent firewall provides a unified security framework for enterprise Large Language Model applications. It addresses prompt injection and jailbreak attacks before model execution, verifies AI-generated responses to reduce hallucinations, protects sensitive enterprise information and strengthens monitoring through risk scoring and audit logging.

5.2 Future Work

- Support multimodal AI systems.
- Introduce adaptive learning for threat detection.
- Deploy in cloud-based enterprise environments.
- Integrate with SIEM platforms.
- Use autonomous AI agents for advanced security decisions.
- Develop a zero-trust AI architecture.

5.3 Professional Learning & Reflection

- Understanding prompt injection and jailbreak threats.
- Learning multi-agent security architecture.
- Understanding hallucination detection and response verification.
- Applying access control, risk scoring and audit logging.
- Developing skills in system architecture, testing, documentation and presentation.

5.4 Individual Reflections

S. Muskan: Focused on evaluation, logging and reporting; learned how audit trails, risk scores and structured testing support enterprise monitoring.

REFERENCES

- [1] C. Zhang, M. Jin, Q. Yu, C. Liu, H. Xue, and X. Jin, “Goal-guided generative prompt injection attack on large language models,” in *Proc. IEEE Int. Conf. Data Mining (ICDM)*, 2025.
- [2] S. S. Ahmed and J. A. Arul Jothi, “Jailbreak attacks on large language models and possible defenses: Present status and future possibilities,” *IEEE*, 2024.
- [3] C. Li, B. Xing, D. Huo, Q. Zhou, Z. Xu, and Y. Wang, “MixHD: A method for detecting hallucinations based on the internal state and output probability of large language models,” in *Proc. IEEE Int. Conf. Acoustics, Speech and Signal Processing (ICASSP)*, 2025.
- [4] H. Xu, W. Zhang, Z. Wang, F. Xiao, R. Zheng, Z. Ba, and K. Ren, “RedAgent: An autonomous agent for context-aware red teaming of LLM jailbreaks,” *IEEE Trans. Dependable Secure Comput.*, vol. 23, no. 3, pp. 6506–6521, 2026, doi: 10.1109/TDSC.2026.3665230.
- [5] F. Xu, H. Hu, C. He, S. Hang, H. Hu, X. Liu, Y. Zhao, Z. Zhou, B. B. Zhu, S.-F. Sun, D. Gu, and S. Wang, “SoK: Robustness in large language models against jailbreak attacks,” in *Proc. IEEE Symp. Security and Privacy*, 2026, pp. 118–137, doi: 10.1109/SP63933.2026.00107.
- [6] P. Lewis *et al.*, “Retrieval-augmented generation for knowledge-intensive NLP tasks,” in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 33, 2020, pp. 9459–9474.
- [7] T. Brown *et al.*, “Language models are few-shot learners,” in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 33, 2020, pp. 1877–1901.
- [8] J. Wei *et al.*, “Chain-of-thought prompting elicits reasoning in large language models,” in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 35, 2022, pp. 24824–24837.
- [9] L. Ouyang *et al.*, “Training language models to follow instructions with human feedback,” in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 35, 2022, pp. 27730–27744.
- [10] N. Shinn *et al.*, “Reflexion: Language agents with verbal reinforcement learning,” in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 36, 2023.

APPENDICES

Appendix A – Proposed System Workflow

The proposed system follows a secure multi-agent workflow for protecting enterprise Large Language Model applications:

User → Authentication & Access Control → LLM Gateway → Multi-Agent Firewall → Enterprise LLM → Hallucination Detection + Sensitive Data Protection → Risk Scoring & Audit Logging → Admin Dashboard / Secure Response

The workflow operates as follows:

1. **User Authentication and Access Control**

The user accesses the enterprise AI application through a secure interface. Authentication verifies the user's identity, while role-based access control restricts functionality according to permissions.

2. **LLM Gateway**

The LLM Gateway securely receives and routes validated user prompts to the multi-agent firewall.

3. **Multi-Agent Firewall**

Every prompt is analyzed before it reaches the Enterprise LLM. The Prompt Injection Detection Agent identifies malicious instructions and hidden attacks, while the Jailbreak Detection Agent identifies attempts to bypass AI safety policies.

4. **Enterprise LLM**

Only verified and approved prompts are forwarded to the Enterprise LLM for response generation.

5. **Hallucination Detection and Sensitive Data Protection**

Generated responses are verified using hallucination detection techniques. Retrieval-Augmented Generation and semantic similarity can be used for factual verification. The Sensitive Data Protection Agent detects confidential enterprise information and masks or blocks sensitive content.

6. **Risk Scoring and Audit Logging**

Each AI interaction is assigned a security risk score. Prompts, responses and firewall decisions are recorded through audit logging.

7. **Admin Dashboard**

The administrator can monitor security events, risk scores, audit information and model evaluation results through the dashboard.

8. Secure Response

After verification and protection, the final response is delivered securely to the authorized user.

Appendix B – Module Summary

Module Name	Inputs	Outputs
Risk Scoring, Audit Logging, Database, Admin Dashboard & Model Evaluation	Interaction metadata, prompts, responses and security decisions	Risk score, audit log, security reports and dashboard information

Module 4 – Risk Scoring, Audit Logging, Database, Admin Dashboard & Model Evaluation

This module assigns security risk scores to interactions, records security events and decisions, stores relevant information in the database and provides administrators with monitoring and evaluation capabilities.

Appendix C – Suggested Test Cases

Case	Input Type	Expected Behavior
TC-01	Normal enterprise query	Request is authenticated, allowed and logged
TC-02	Prompt injection attempt	Attack is detected and prompt is blocked or sanitized
TC-03	Jailbreak attempt	Jailbreak behavior is detected and request is blocked
TC-04	Potentially hallucinated / unsupported response	Response is flagged for hallucination verification
TC-05	Confidential enterprise data	Sensitive information is masked or blocked
TC-06	High-risk interaction	Elevated risk score is generated and an audit event is recorded
TC-07	Unauthorized user request	Access is denied according to role permissions
TC-08	Safe verified response	Response passes security checks and is delivered to the user

The test cases should be executed using controlled inputs and the actual implementation results should be recorded in the final evaluation.

Appendix D – Implementation Evidence

The implementation section of the presentation indicates that implementation was performed, but the extracted PPT content does not provide sufficient details of the actual source code, numerical evaluation results or complete screenshots. Therefore, the following evidence should be inserted from the actual project implementation:

1. User interface screenshot.
2. Authentication and role-based access-control screen.
3. LLM Gateway workflow.
4. Multi-Agent Firewall implementation.
5. Prompt Injection Detection Agent output.
6. Jailbreak Detection Agent output.
7. Hallucination Detection Agent output.
8. Sensitive Data Protection output.
9. Risk scoring result.
10. Audit logging/database output.
11. Admin dashboard screenshot.
12. Model evaluation results.
13. Sample secure and blocked responses.
14. Source-code excerpts for the major modules.




Fig. 4.1 Secure Chat Interface Showing Prompt Security Detection




Fig. 4.2 Admin Dashboard Showing Model Evaluation and Audit Logs




Fig. 4.3 Request Logs Showing Security Actions and Risk Scores

Appendix E – Ethics / Institutional Considerations

The proposed system should be developed and evaluated responsibly because it handles enterprise AI interactions and may process confidential information.

The following ethical and security considerations apply:

1. Protection of Confidential Information

Enterprise data used for testing should be controlled and sensitive information should be masked or removed where appropriate.

2. Responsible AI Usage

The limitations of Large Language Models, including the possibility of hallucinated or misleading responses, should be clearly recognized.

3. Transparent Reporting

Results should be reported accurately without fabricating security or model-performance results.

4. Access Control

Only authorized users should be permitted to access enterprise AI functions and security information.

5. Audit and Accountability

Security events, prompts, responses and firewall decisions should be logged appropriately to support monitoring and compliance.

6. **AI-Assisted Development Disclosure**

Any AI-assisted development, coding or content-generation tools used during the project should be disclosed where required by institutional policy.

7. **Secure Data Handling**

Audit logs and databases should be protected against unauthorized access or modification.

8. **Human Oversight**

High-risk security decisions should be monitored appropriately rather than relying entirely on autonomous AI decisions.

Appendix F – Future Enhancement Areas

The presentation identifies the following future enhancements:

- Extension to **multimodal AI systems**.
- **Adaptive learning** for improved threat detection.
- **Cloud deployment** for large-scale enterprise environments.
- **SIEM integration** for stronger organizational cybersecurity operations.
- **Autonomous AI agents** for advanced security decisions.
- **Zero-trust AI architecture** for enhanced enterprise protection.

These enhancements can improve scalability, adaptability and security for enterprise Generative AI applications.